

# Continuous Delivery with Docker

Tobias Schwab

# Myself

- Tobias Schwab
- tobias.schwab@dynport.de
- [www.dynport.de](http://www.dynport.de)
- [twitter.com/tobstarr](https://twitter.com/tobstarr)
- [github.com/tobstarr](https://github.com/tobstarr)

# Philosophie

- continuous delivery: deploy multiple times a day
- canary releases
- “never touch a running system”
- “Immutable Infrastructure and Disposable Components”
- don't fix it, if it can be replaced

# Theory

- AWS
- AMI based deployments
- Elastic Load Balancer
- AutoScaling Groups
- S3, RDS, ...

# Reality

- privacy concerns: AWS not an option
- hoster we could not pick
- first no, then proprietary and unreliable API
- flash based infrastructure management
- limited capacity
- we were the biggest customer

# Build

- manual
  - start and attach container
  - install required packages
  - checkout application code
  - run build management tool
  - bad: not reproducible
  - bad: does not utilise caching

# Build

- chef/puppet/...
  - start an attach container
  - run chef/puppet/... client
  - good: automated and documented
  - bad: does not utilise caching

# Dockerfile

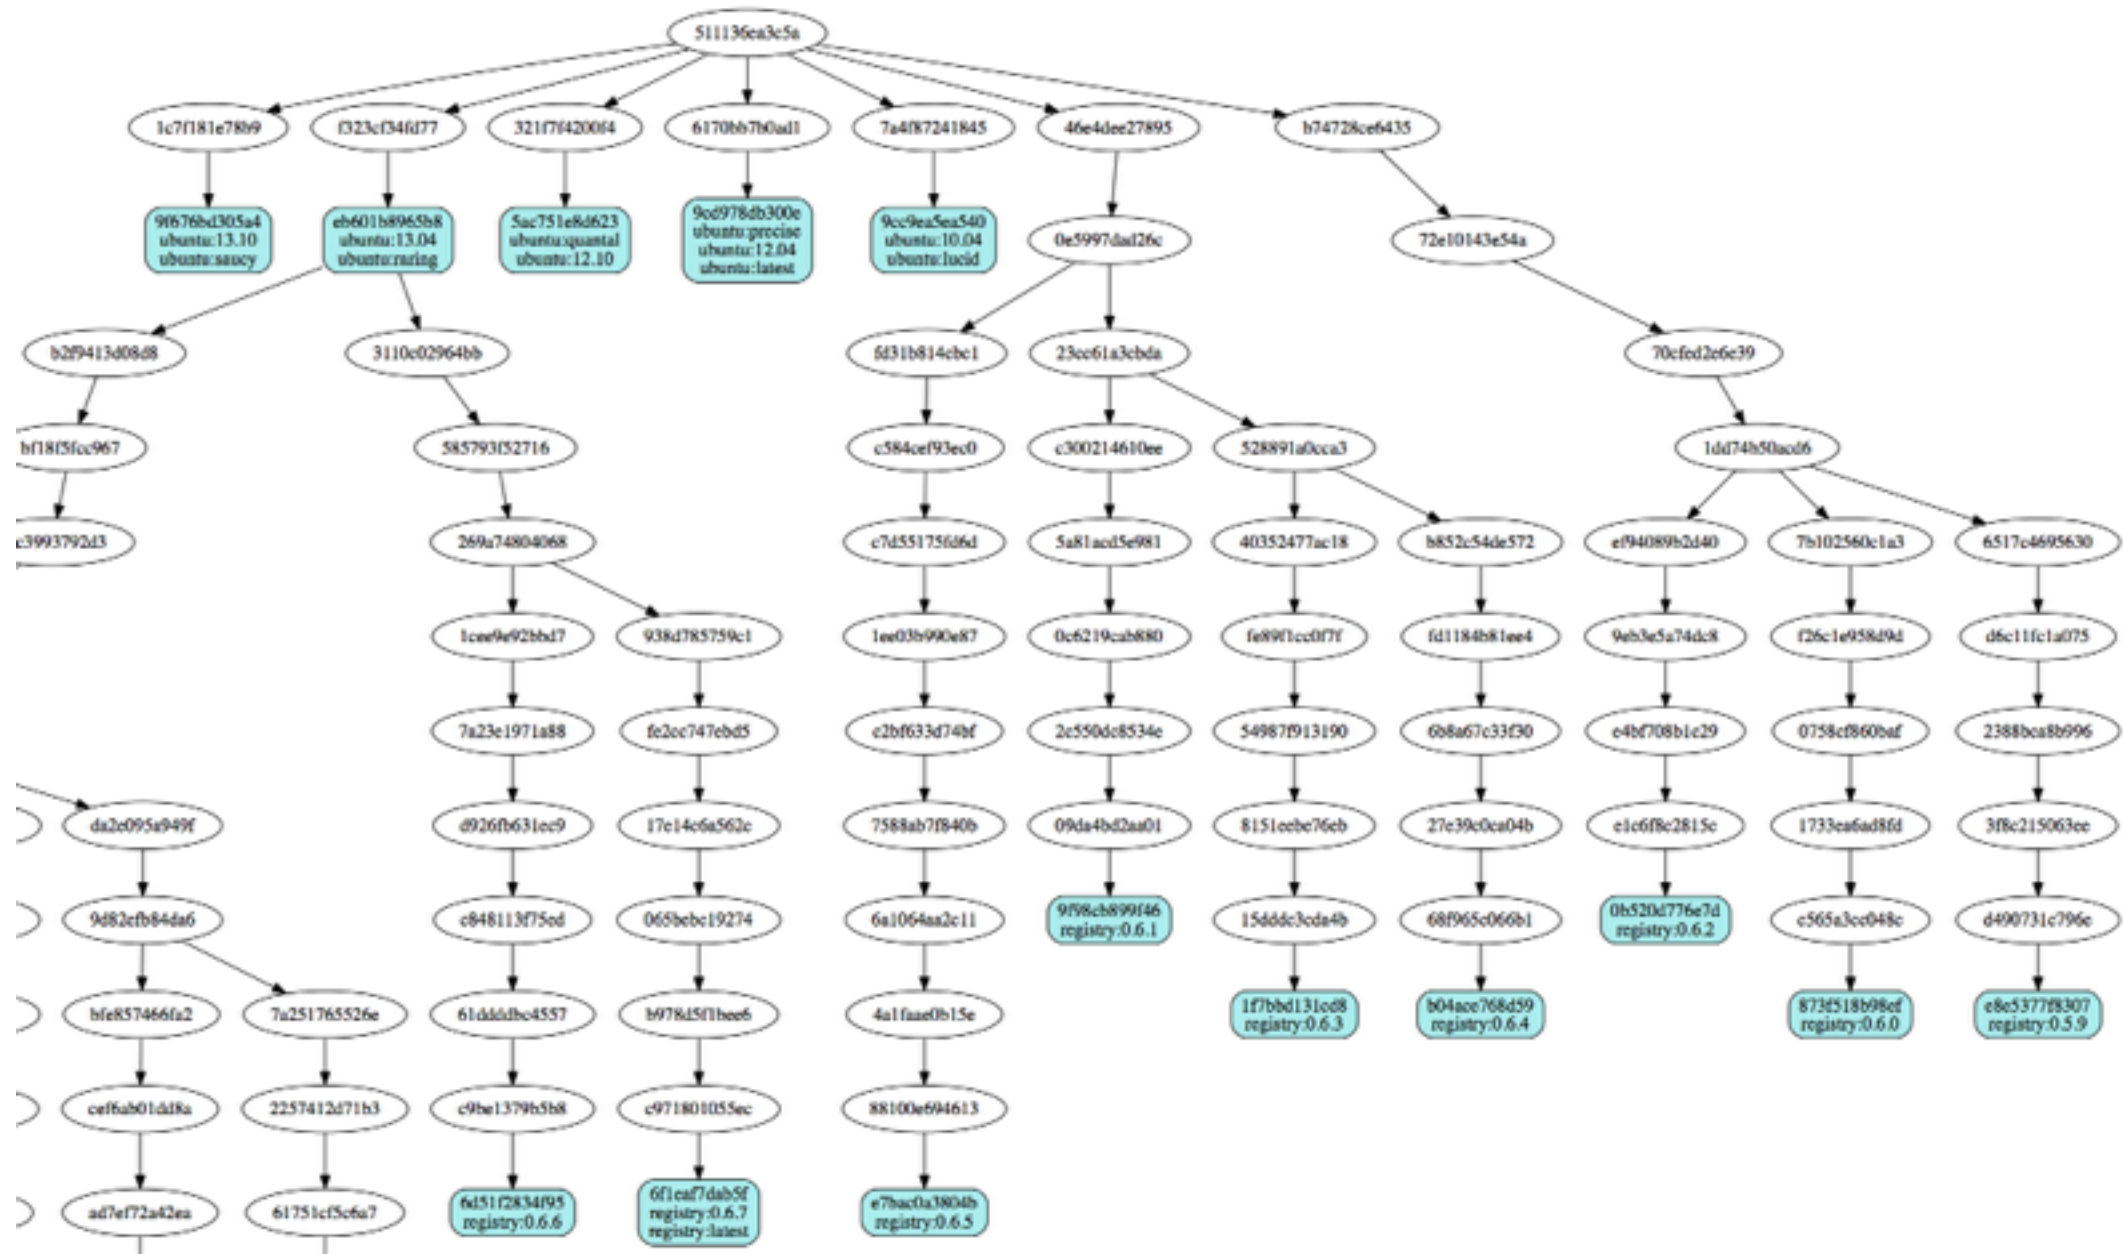
- simple, plain text script to create images
- commands:
  - FROM: base image to use
  - RUN: execute shell command
  - ENV: set environment variable
  - ADD: write local file to image
  - ENTRYPOINT: start command for containers
- others: MAINTAINER, EXPOSE, CMD, USER, VOLUME, WORKDIR, ONBUILD



# Caching

- statement based: each step creates a new image
- existing steps (command tree exists) are re-used
- tricky: “non functional” commands (e.g. apt-get update/upgrade)
- use ENV or comments to break caching of non functional commands

# Caching



# Configuration Management

- “store config in the environment” (<http://12factor.net/config>)
- dependency injected with start of container
- same image for
  - development
  - testing
  - staging
  - production

# Don'ts

- full blown VMs
- ssh daemon inside containers
- syslog daemon inside containers (sometimes needed)
- user management: everything can run as root
- chef/puppet/... => makes caching useless

# Build Management Tools

- candidates: bundler, pip, mvn, carton, composer, ...
- problem with caching: bmt are slow when started with “clean slate”
- option 1: add bmt manifest before code
  - bmt needs to run only when manifest changes
- option 2: use pre-bundled base images
  - bmt only needs to work the delta
  - re-build base images from time to time
- option 3: combine option 1 and option 2

# Use Case: Ruby on Rails

- Problems
  - unicorn: Rack HTTP server for fast clients
  - static assets
  - logging: default ruby syslog library uses syscall (needs local syslog daemon)
- Solution
  - run 3 daemons in 1 container: unicorn, nginx and rsyslogd
  - upstart
    - ENTRYPOINT ["/sbin/init"]
    - load ENV from /proc/1/environ
  - foreman

# Multi-Host

- image distribution via docker registry
- weighted load balancing via HAProxy
- SSL termination via nginx in front of HAProxy

# Registry

- push and pull images
- public
- private
- backends: local, S3, Elastics, Google Cloud Storage, hosted



# Load Balancing

- HAProxy
- pool configuration stored in redis/etcd
- config update
  - compile config files from stored configuration
  - upload via ssh
  - verify on remote hosts
  - replace current config with verified one
  - reload

# HAProxy

```
1 global
2   daemon
3   pidfile /var/run/haproxy.pid
4   stats socket /tmp/haproxy.socket level admin
5   log 192.168.0.4 user info
6
7 defaults
8   mode http
9   maxconn 50000
10  timeout connect 5000ms
11  timeout client 50000ms
12  timeout server 50000ms
13  option http-server-close
14  option http-pretend-keepalive
15  option httplog
16  log global
17
18 frontend in-stats
19   bind *:20002
20   default_backend stats
21
22 backend stats
23   stats uri /_stats
24   stats refresh 5
25
26 frontend in-rails
27   bind 192.168.0.6:8080
28   default_backend rails
29
30 backend rails
31   rspadd X-Proxy:\ cnc-5f22a980
32   option httpchk GET /_status
33
34   # server cnc-369bf803:f753afc24d79:192c3f992501
35   acl cnc-369bf803:f753afc24d79:192c3f992501 srv_id 1
36   rspadd X-Host:\ cnc-369bf803:f753afc24d79:192c3f992501 if cnc-369bf803:f753afc24d79:192c3f992501
37   server cnc-369bf803:f753afc24d79:192c3f992501 192.168.0.19:49382 check inter 5s id 1 weight 100
38
39   # server cnc-369bf803:f753afc24d79:5dd19b1da07a
40   acl cnc-369bf803:f753afc24d79:5dd19b1da07a srv_id 2
41   rspadd X-Host:\ cnc-369bf803:f753afc24d79:5dd19b1da07a if cnc-369bf803:f753afc24d79:5dd19b1da07a
42   server cnc-369bf803:f753afc24d79:5dd19b1da07a 192.168.0.19:49383 check inter 5s id 2 weight 100
```

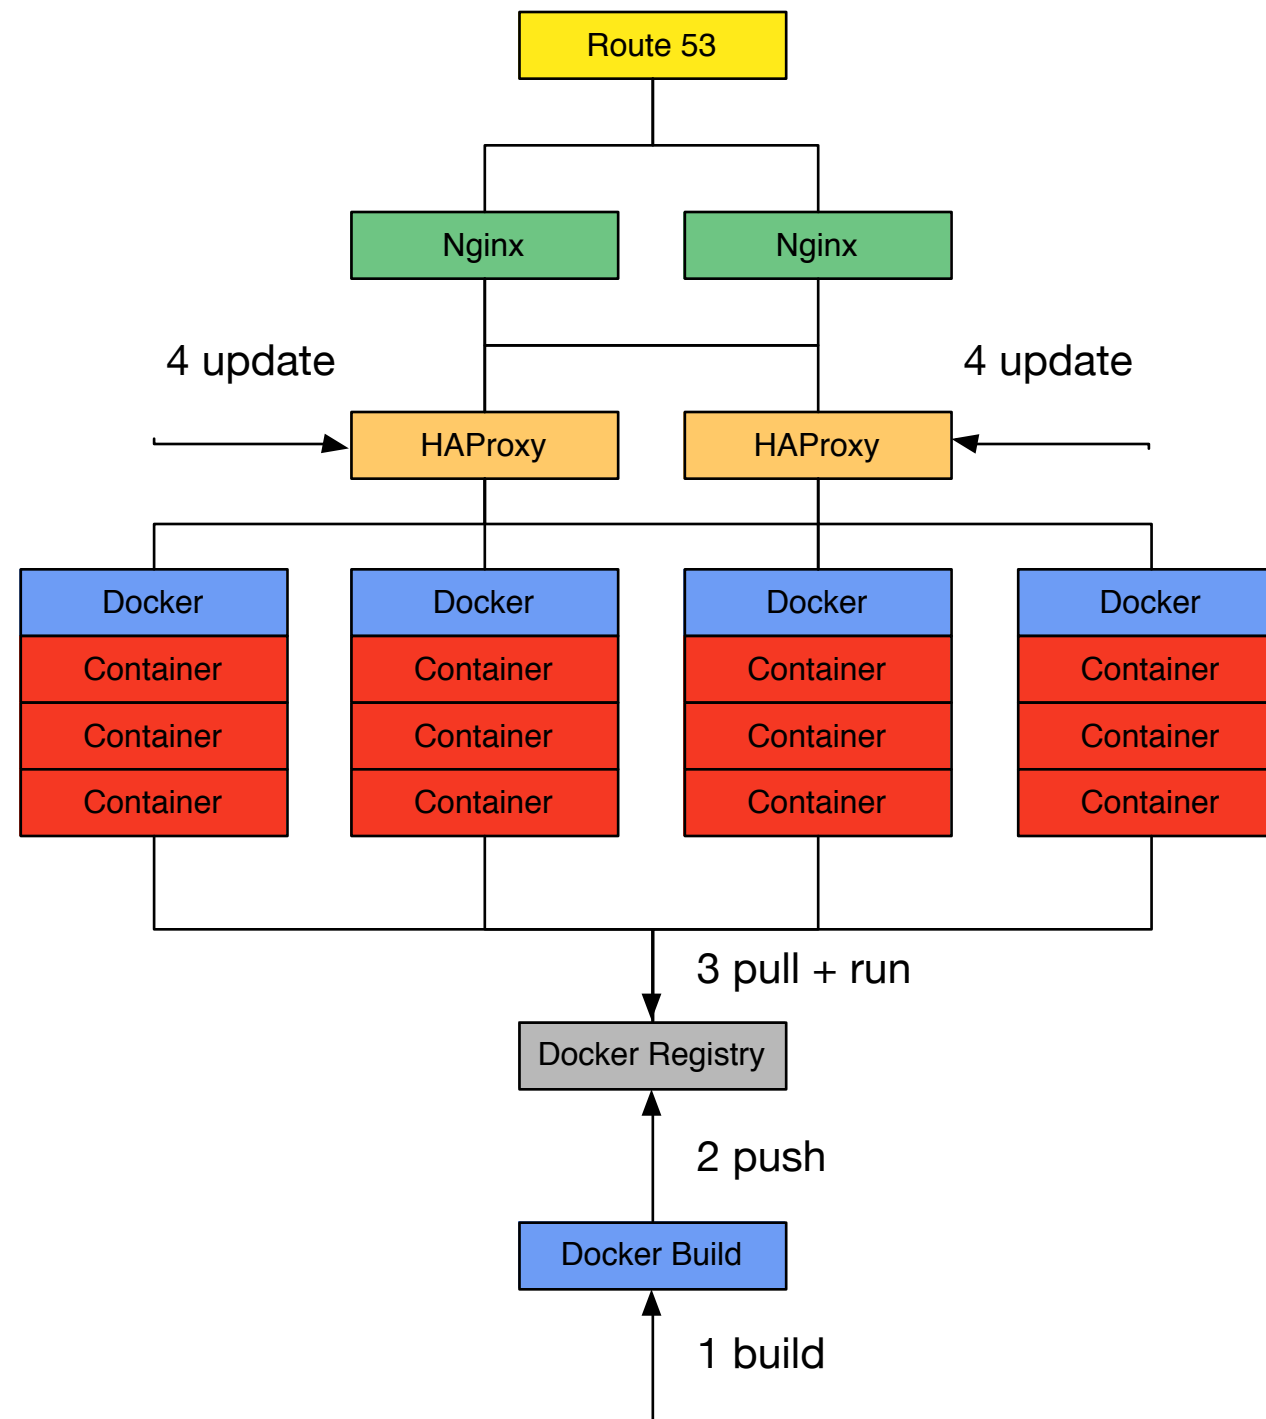
# HAProxy

	Queue			Session rate			Sessions					Bytes		Denied		Errors			Warnings		Status
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	LbTot	In	Out	Req	Resp	Req	Conn	Resp	Retr	Redis	
cnc-535934d7:b2b46b4425a7:4074cf7ba135	0	0	-	2	6		0	29	-	1 215	1 215	38 759 651	26 592 153		0		0	0	0	0	3m28s UP
cnc-82f8a2c8:b2b46b4425a7:b2674302a38c	0	0	-	2	6		1	50	-	1 170	1 170	28 607 397	29 647 058		0		0	0	0	0	1m10s UP
cnc-87c73527:b2b46b4425a7:e249016faa9a	0	0	-	2	8		0	8	-	1 283	1 283	9 625 072	24 274 168		0		0	0	0	0	12m50s UP
cnc-c1b3dd12:b2b46b4425a7:26e76d513ce4	0	0	-	2	8		8	31	-	1 257	1 257	17 259 361	33 735 903		0		0	0	0	0	2m15s UP
Backend	0	0		9	22		9	59	0	4 925	4 925	94 251 481	114 249 282	0	0		0	0	0	0	12m50s UP

# Deployment Pipeline

- commit triggers new image build
- build suite executed with image
- image is pushed to registry if tests passed
- optional: start image with staging ENV settings for manual testing
- start image with production ENV for last pre-flight tests
- deploy image to more hosts
- update load balancer (canary or green/blue)
- monitor new containers/image

# Deployment Pipeline



# Logging

- NOT inside containers
- remote syslog (when possible)
- alternative: local syslog relay inside container

```
2013-11-20T16:00:00.003829+00:00 1d63a4fc82a1 nginx: [2013-11-20 16:00:00.003829+00:00] method=GET status=200 length=38173 pid=27771 rev=5f3cf70cc90c image_id=b8d7736b9986 uuid=da7d5065-a544-4c6a-88dd-7c0dbc57fdb6 action=home#index etag=2dffe9baf37402334042dde91c055db5 rack=0.140591 db=0.035836/12 db_cache=0.002128/3 total=0.146 upstream_time=0.143 ua="Mozilla/5.0 (Linux; Android 4.2.2; de-de; SAMSUNG GT-I9195 Build/JDQ39) AppleWebKit/535.19 (KHTML, like Gecko) Version/1.0 Chrome/18.0.1025.308 Mobile Safari/535.19" uri="/" ref="-"
```

- host: docker host, container\_id
- code: image\_id, revision
- request: request\_id, action, status\_code, etag, times, calls

# Metrics

- OpenTSDB
  - “distributed, scalable Time Series Database”
  - HBase
  - Tags / Dimensions
- from syslog via udp (StatsD “like”)
- rickshaw.js for graphs
- compare status codes, counts and times between actions of two revisions




# OpenTSDB

```
nginx.timers.total.count 1389589200 1 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589201 2 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589202 3 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589204 4 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589208 6 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589209 7 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589210 9 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589212 11 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589213 14 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589215 15 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589216 18 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589217 20 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589218 21 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589221 24 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589224 27 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589225 28 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589226 32 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589227 33 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589228 35 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589229 40 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589230 64 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589231 129 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589232 149 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589233 157 revision=bb103bbf37f3 image_id=a0c79553f35f
nginx.timers.total.count 1389589234 159 revision=bb103bbf37f3 image_id=a0c79553f35f
```

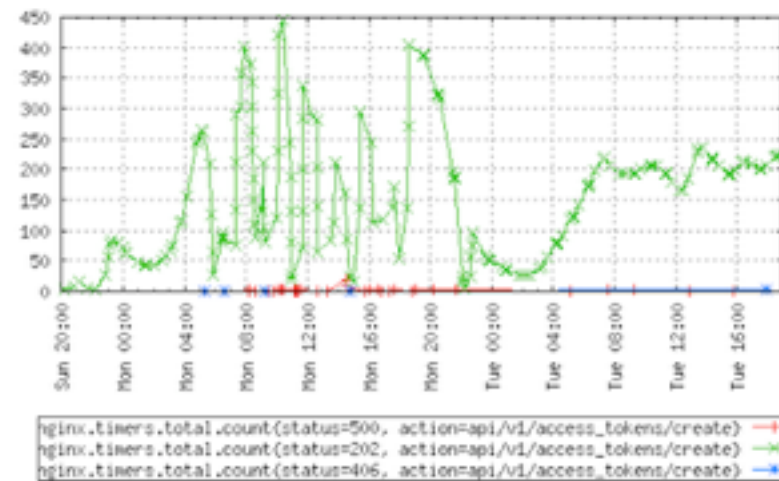


# Metrics

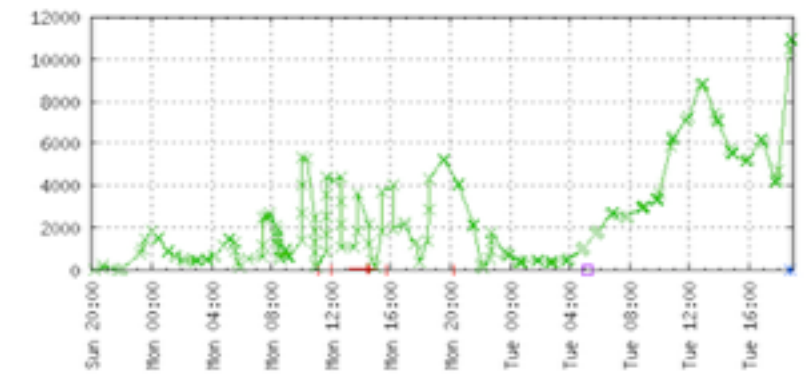
api/v1/photos/create




api/v1/access\_tokens/create



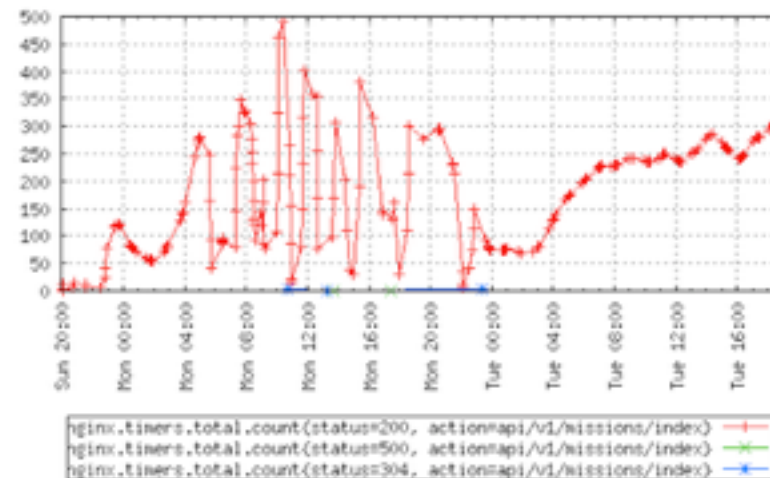
api/v1/ratings/update



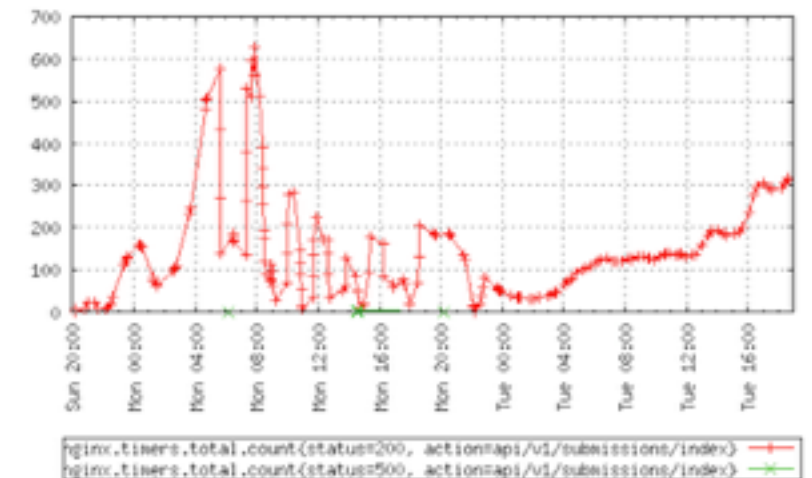
api/v1/photos/index




api/v1/missions/index



api/v1/submissions/index

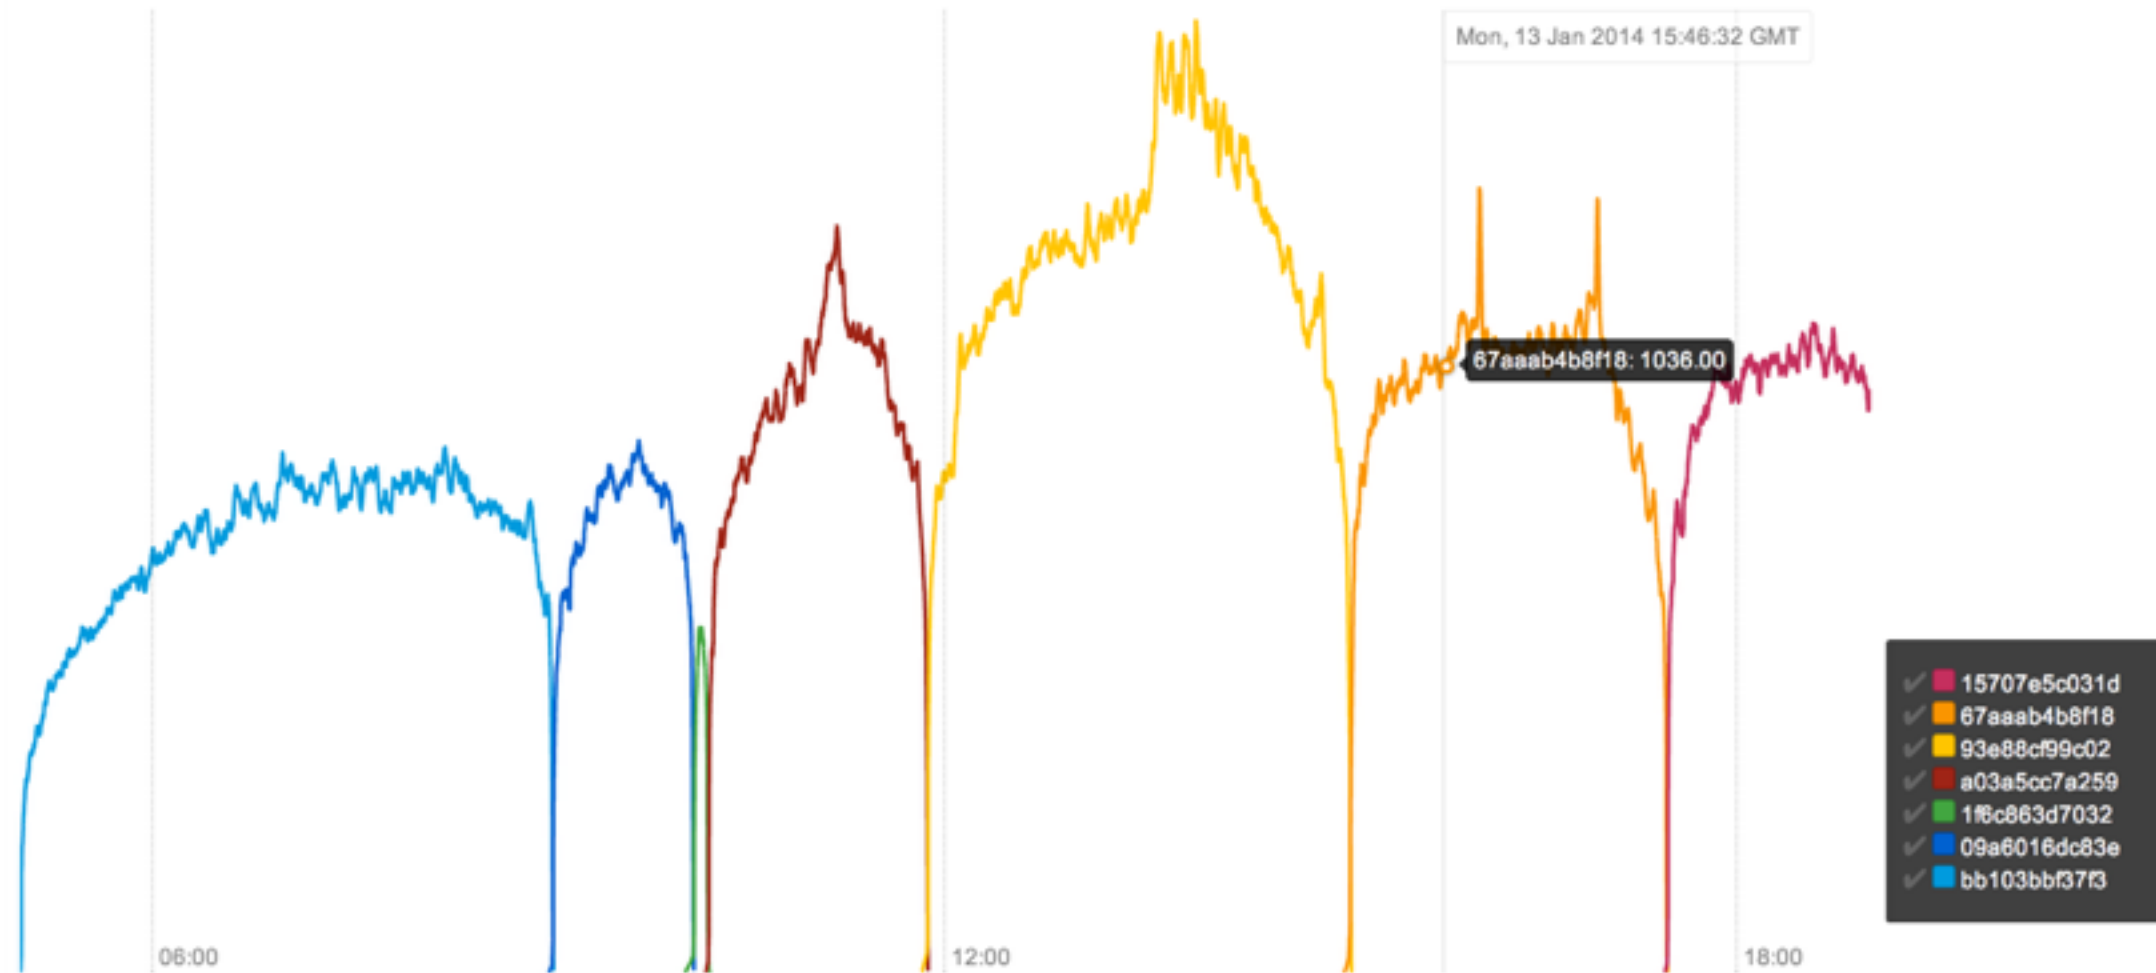


# Metrics



request counts by revision


# Metrics



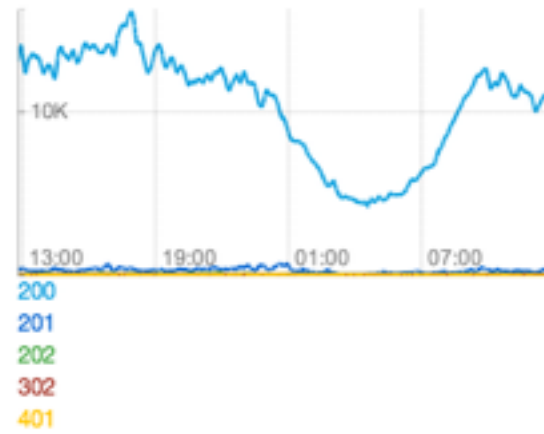
# Metrics

## Requests


### Action




### Status Total



### Total 95%




### DB 95%



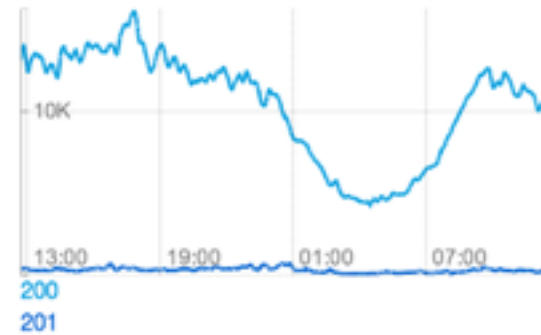
### Rack 95%



### 3xx



### 2xx




### 5xx




# Metrics


Requests




Status Total




Load




Memory Used



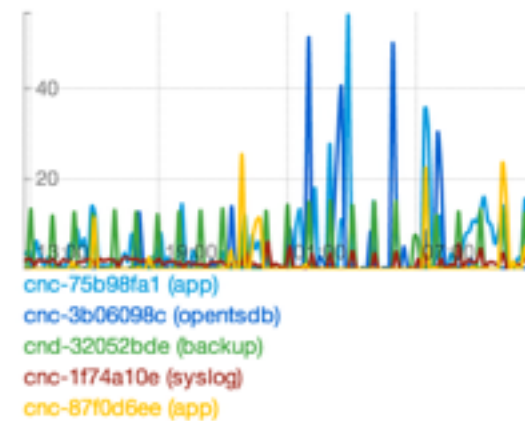
Swap Used




Disk Use



IOWait



User



System

Disk IO

Disk Read

# Docker reduces

- external dependencies (“rubygems/github slow/unreliable/down”) after image is built
- “did work on my machine/staging”: same OS package versions, configuration and code in all stages
- unused CPU cycles
- number of hosts
- feedback times
- time to get new host online
- bottlenecks: hosts are more flexible

# VS. AWS

- HAProxy much more flexible
  - multiple containers per host
  - balancing weights
- faster build process
- faster deployments
- instance flexibility

# Resources

- [docker.io](https://docker.io)
- [opentsdb.net](https://opentsdb.net)
- [haproxy.1wt.eu](https://haproxy.1wt.eu)
- [continuousdelivery.com](https://continuousdelivery.com)
- [chadfowler.com/blog/2013/06/23/immutable-deployments/](https://chadfowler.com/blog/2013/06/23/immutable-deployments/)
- [12factor.net](https://12factor.net)



Questions?!?

Thank you!